

Core I—Introduction to HPC

Session V: OpenMP (1/2)

Dr. Weinzierl

Michaelmas term 2019

Outline



OpenMP
Thread communication

OpenMP is a ...


- ▶ abbreviation for *Open Multi Processing*
- ▶ allows programmers to annotate their C/FORTRAN code with parallelism specs
- ▶ portability stems from compiler support
- ▶ standard defined by a consortium (www.openmp.org)
- ▶ driven by AMD, IBM, Intel, Cray, HP, Nvidia, ...
- ▶ old thing currently facing its fourth generation (1st: 1997, 2nd: 2000; 3rd: 2008; 4th: 2013)

A first example

```
const int size = ...  
int a[size];  
#pragma omp parallel for  
for( int i=0; i<size; i++ ) {  
    a[i] = i;  
}
```

- ▶ OpenMP for C is a preprocessor(pragma)-based extension (annotations, not API)
 - ▶ Implementation is up to the compiler (built into recent GNU and Intel compilers; before additional precompiler required)
 - ▶ Implementations internally rely on libraries (such as pthreads)
 - ▶ Annotations that are not understood should be ignored (don't break old code)
- ▶ Syntax conventions
 - ▶ OpenMP statements always start with `#pragma omp`
 - ▶ Before GCC 4, source-to-source compiler replaced only those lines
- ▶ OpenMP originally written for BSP/Fork-Join applications
- ▶ OpenMP usually abstracts from machine characteristics
(number of cores, threads, ...)

A first example



```
const int size = ...  
int a[size];  
#pragma omp parallel for  
for( int i=0; i<size; i++ ) {  
    a[i] = i;  
}
```

Compilation and execution:

- ▶ GCC: `-fopenmp`
- ▶ Intel: `-openmp`
- ▶ Some systems require `#include <omp.h>`
- ▶ Set threads: `export OMP_NUM_THREADS=2`

What happens—usage

```
>icc -fopenmp test.cpp  
test.cpp(22): (col. 1) remark: \  
OpenMP DEFINED REGION WAS PARALLELIZED.
```

```
>export OMP_NUM_THREADS=4  
>./a.out
```

```
const int size = ...  
int a[size];  
#pragma omp parallel for  
for( int i=0; i<size; i++ ) {  
    a[i] = i;  
}
```

- ▶ Code runs serially until it hits the pragma
- ▶ System splits up for loop into chunks (we do not yet know how many)
- ▶ Chunks then are deployed among the four threads
- ▶ All threads wait until loop has terminated on all threads, i.e. it synchronises the threads \Rightarrow bulk synchronous processing (bsp)
- ▶ Individual threads may execute different instructions from the loop concurrently (designed for MIMD machines)

OpenMP execution model

Explicit scoping:

```
#pragma omp parallel for
{
    for( int i=0; i<size; i++ ) {
        a[i] = i;
    }
}
```

Implicit scoping:

```
#pragma omp parallel for
for( int i=0; i<size; i++ ) {
    a[i] = i;
}
```

- ▶ Master thread vs. worker threads
- ▶ Fork/join execution model with implicit synchronisation (barrier) at end of scope
- ▶ Nested parallel loops possible (though sometimes very expensive)
- ▶ Shared memory paradigm (everybody may access everything)

Some OMP functions

```
int numberOfThreads = omp_get_num_procs();  
  
#pragma omp parallel for  
for( int i=0; i<size; i++ ) {  
  
    int thisLineCodelsRunningOn = omp_get_thread_num();  
  
}
```

- ▶ No explicit initialisation OpenMP required in source code
- ▶ Abstract from threads—setting thread count is done by OS
- ▶ Error handling (to a greater extent) not specified by standard
- ▶ Functions almost never required (perhaps for debugging)

Parallel loops in action


```
#pragma omp parallel
{
    for( int i=0; i<size; i++ ) {
        a[i] = a[i]*2;
    }
}
```

```
#pragma omp parallel for
{
    for( int i=0; i<size; i++ ) {
        a[i] = a[i]*2;
    }
}
```

Observations:

- ▶ Global loop count is either *size* or *threads*·*size*
- ▶ We run into race conditions
- ▶ These result from dependencies (read-write, write-read, write-write) on *a[i]*

Parallel loops and BSP



```
#pragma omp parallel
{
    #pragma omp for
    for( int i=0; i<size; i++ ) {
        a[i] = a[i]*2;
    }
}
```


- ▶ `omp parallel` triggers fork technically, i.e. spawns the threads
- ▶ `for` decomposes the iteration range (logical fork part)
- ▶ `omp parallel for` is a shortcut
- ▶ BSP's join/barrier is done implicitly at end of the parallel section

Requirements for parallel loops

```
#pragma omp parallel for
{
    for( int i=0; i<size; i++ ) {
        a[i] = a[i]*2;
    }
}
```

- ▶ Loop has to follow plain initialisation-condition-increment pattern:
 - ▶ Only integer counters
 - ▶ Only plain comparisons
 - ▶ Only increment and decrement (no multiplication or any arithmetics)
- ▶ Loop has to be countable (otherwise splitting is doomed to fail).
- ▶ Loop has to follow single-entry/single-exit pattern.

Data consistency in OpenMP's BSP model



```
#pragma omp parallel
{
  #pragma omp for
  for( int i=0; i<size; i++ ) {
    a[i] = a[i]*2;
  }
}
```

- ▶ No assumptions which statements run technically concurrent
 - ▶ Shared memory without any consistency model
 - ▶ No inter-thread communication (so far)
- ⇒ Data consistency is developer's responsibility

Concept of building block: OpenMP Introduction

- ▶ Content
 - ▶ OpenMP syntax basics
 - ▶ OpenMP runtime model
 - ▶ OpenMP functions
- ▶ Expected Learning Outcomes
 - ▶ The student can *translate* and use an application with OpenMP support
 - ▶ The student can *explain* with the OpenMP execution model

Outline



OpenMP
Thread communication

Thread communication

We distinguish two different communication types:

- ▶ Communication through the join
- ▶ Communication inside the BSP part (not “academic” BSP)

Critical section: Part or code that is ran by at most one thread at a time.

Reduction: Join variant, where all the threads reduce a value into one single value.

- ⇒ Reduction maps a vector of $(x_0, x_1, x_2, x_3), \dots, x_{p-1}$ onto one value x , i.e. we have a all-to-one data flow
- ⇒ Inter-thread communication realised by data exchange through shared memory
- ⇒ Fork can be read as one-to-all information propagation (done implicitly by shared memory)

Critical sections

```
#pragma omp critical (mycriticalsection)
{
    x *= 2;
}
...
#pragma omp critical (anothersection)
{
    x *= 2;
}
...
#pragma omp critical (mycriticalsection)
{
    x /= 2;
}
```

- ▶ Name is optional (default name i.e. does not block with other sections with a name)
- ▶ Single point of exit policy \Rightarrow return, break, ... not allowed within critical section
- ▶ For operations on built-in/primitive data types, an *atomic operation* is usually the better, i.e. faster choice (cmp. CCS module)

Recap: Requirements for parallel loops

```
#pragma omp parallel for
{
    for( int i=0; i<size; i++ ) {
        a[i] = a[i]*2;
    }
}
```

- ▶ Loop has to follow plain initialisation-condition-increment pattern:
 - ▶ Only integer counters
 - ▶ Only plain comparisons
 - ▶ Only increment and decrement (no multiplication or any arithmetics)
- ▶ Loop has to be countable (otherwise splitting is doomed to fail).
- ▶ Loop has to follow single-entry/single-exit pattern.
- ▶ Loop copies all share the memory.

Consistency observation:

- ▶ All attributes are shared
 - ▶ Besides the actual loop counter (otherwise splitting wouldn't work)
- ⇒ There has to be support for non-shared data

The shared default clause

```
double result = 0;
#pragma omp parallel for
for( int i=0; i<size; i++ ) {
    result += a[i] * b[i]; // race, but doesn't bother us here
}
```

```
double result = 0;
#pragma omp parallel for shared(result)
for( int i=0; i<size; i++ ) {
    result += a[i] * b[i];
}
```

- ▶ By default, all OpenMP threads share all variables, i.e. variables declared outside are visible to threads
- ▶ This sharing can be made explicit through the clause `shared`
- ▶ Explicit shared annotation is kind of good style (improved readability)

Thread-local variables

```
double result = 0;
#pragma omp parallel for
for( int i=0; i<size; i++ ) {
    double result = 0.0;
    result += a[i] * b[i];
}
```

Without OpenMP pragma:

- ▶ C/C++ allows us to “redefine” variable in inner scope
- ▶ Hides/shadows outer `result`
- ▶ We may not forget the second initialisation; otherwise garbage

With OpenMP pragma:

- ▶ OpenMP introduces (concurrent) scope of its own
- ▶ Scope-local variables are thread-local variables
- ▶ These are *not* shared

shared vs. private clauses

```
double result = 0;
#pragma omp parallel for private(result)
for( int i=0; i<size; i++ ) {
    result += a[i] * b[i];
}
```

- ▶ `private` is the counterpart of `shared`, i.e. each thread works on its own copy
- ▶ Basically, the copying is similar to the following fragment:

```
double result = 0;
#pragma omp parallel
{
    double result;
    #pragma omp for
    for( int i=0; i<size; i++ ) {
        result += a[i] * b[i];
    }
}
```

- ⇒ In this example, `result` within thread is not initialised (garbage)!
- ⇒ In this example, data within `result` is lost throughout join!

Example

```
int x;  
x=40;  
#pragma omp parallel for private (x)  
for (int i=0; i<10; i++) {  
    x = i; // try to remove this line  
    std::cout << "x=" << x << "_on_thread_"  
                << omp_get_thread_num() << std::endl;  
}  
std::cout << "x=" << x << std::endl;
```

Example

```
int x;  
x=40;  
#pragma omp parallel for private (x)  
for (int i=0; i<10; i++) {  
    x = i; // try to remove this line  
    std::cout << "x=" << x << "_on_thread_"  
               << omp_get_thread_num() << std::endl;  
}  
std::cout << "x=" << x << std::endl;
```

Observations:

- ▶ If we comment out `x=i`, `x` is not properly initialised.
- ▶ `x` in the last line always equals 40.

Now remove initialisation and write `firstprivate(x)`.

Copy policies

- ▶ `default` Specifies default visibility of variables
- ▶ `firstprivate` Variable is private, but initialised with value from surrounding
- ▶ `lastprivate` Variable is private, but value of very last iteration is copied over to outer variable

Case study: scalar product

Serial starting point:

```
double result = 0;
for( int i=0; i<size; i++ ) {
    result += a[i] * b[i];
}
```

Parallel variant:

```
double result = 0.0;
#pragma omp parallel
{
    double myResult = 0.0;
    #pragma omp for
    for( int i=0; i<size; i++ ) {
        myResult += a[i] * b[i];
    }
    #pragma omp critical
    result += myResult;
}
```

Observations:

- ▶ Avoid excessive synchronisation
- ▶ Type of operation is called *reduction* (as defined before)
- ▶ We may not use `result` to accumulate because of races
- ▶ We may not hide `result` as we then lose access to outer variable

Concept of building block

- ▶ Content
 - ▶ Introduce three types of data flow/usage of shared memory: making memory available to all threads throughout fork, sharing data throughout computation, reducing data at termination
 - ▶ Introduce private variables
 - ▶ Study semantics and usage of critical sections
- ▶ Expected Learning Outcomes
 - ▶ The student *can use* critical sections
 - ▶ The student *knows difference* between private and shared variables
 - ▶ The student can *identify race conditions* and *resolve* them through critical sections for given code snippet